Chosen 02 | Observed Latency vs Mixed Rank

All three datasets | Zero-shot + Few-shot | original_glm

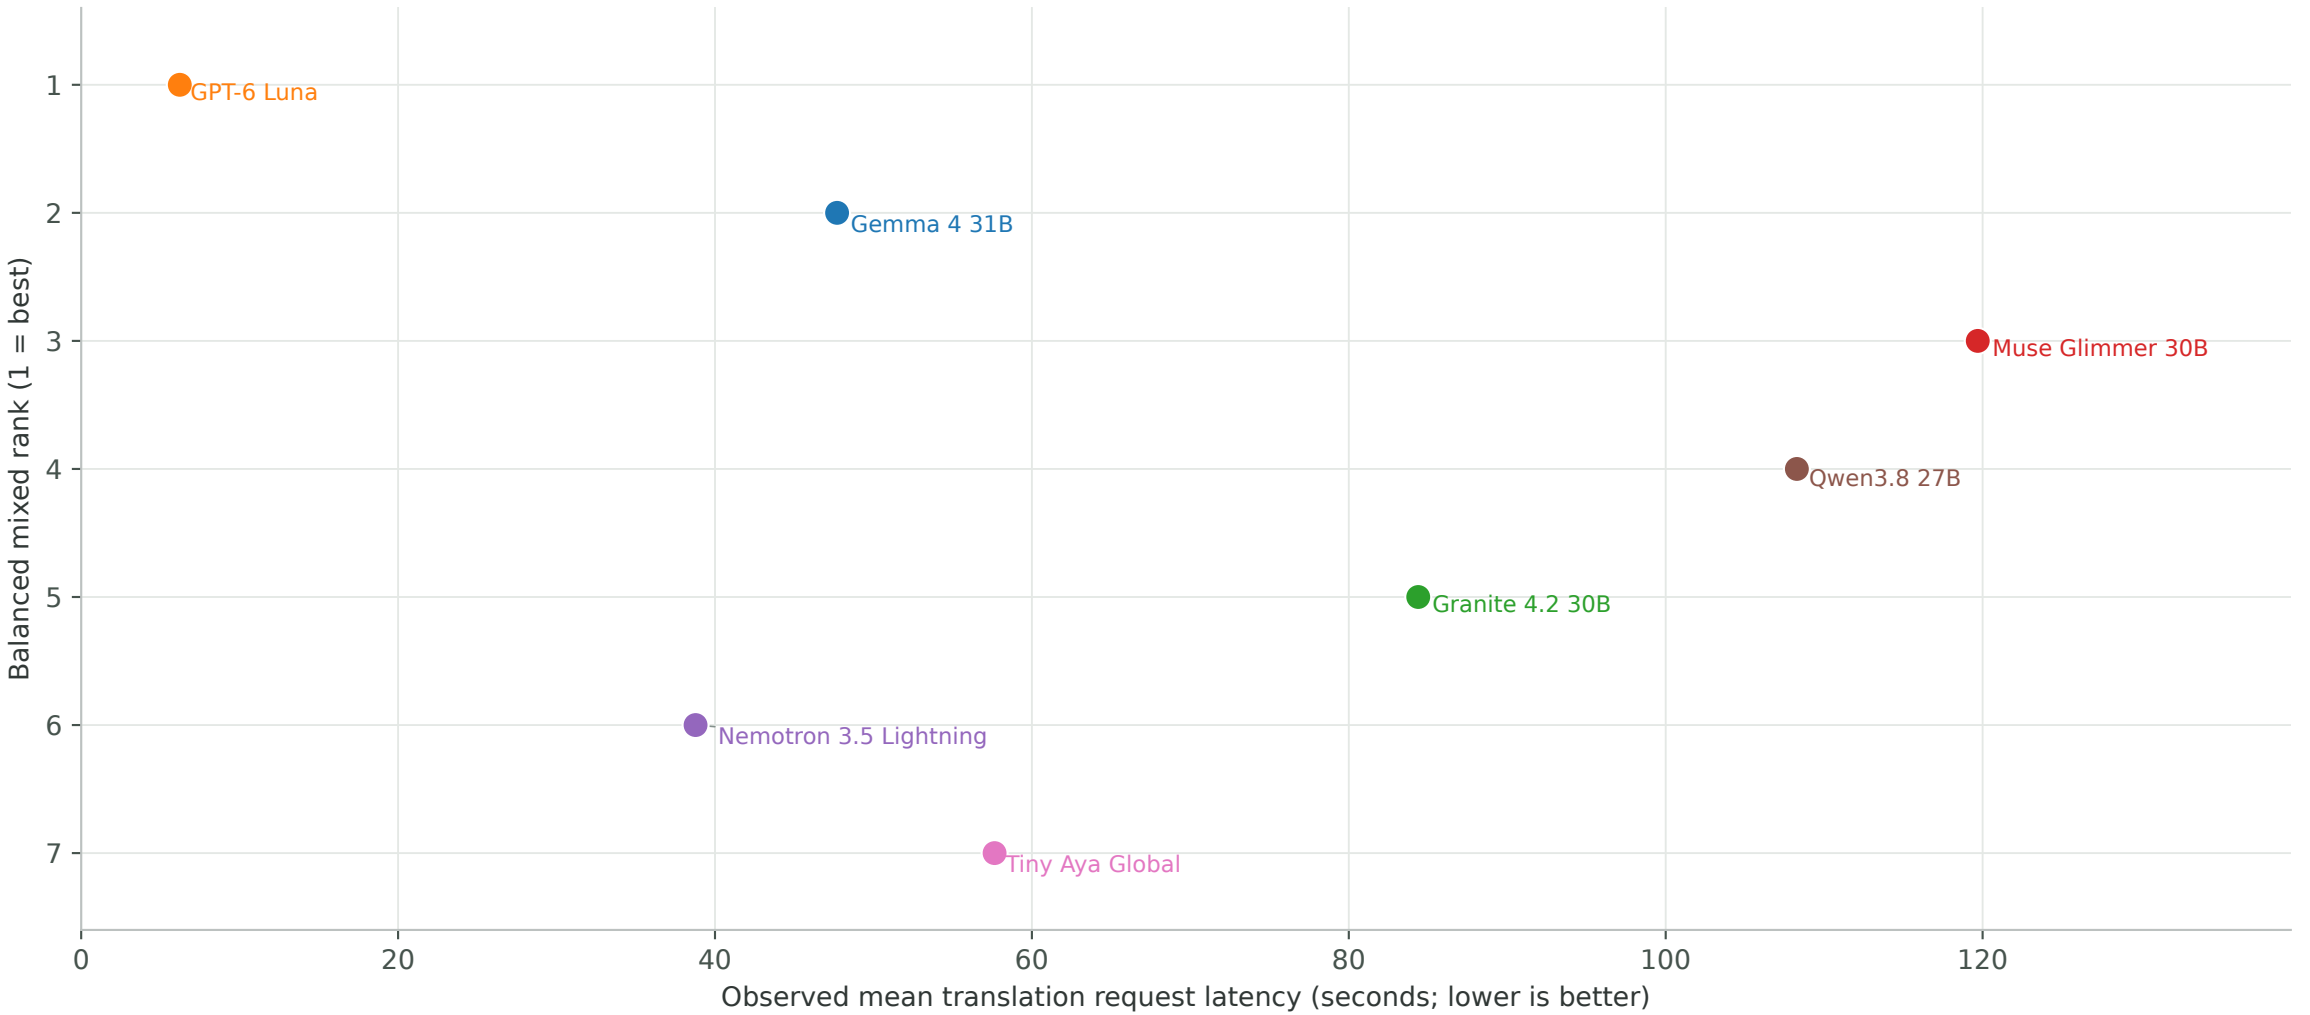

Upper-left is preferable. X: recorded generation latency, not evaluator time. Overall means weight the six dataset/prompt strata equally.
Observed deployments differ in hardware, serving/API setup, concurrency and workloads; this is not a hardware-controlled speed comparison.
Timing uses all recorded timed requests; quality uses metric-specific scored support and is not proven row-matched to timing. No monetary cost inferred.