Direction-Level Performance Distributions | Original GLM | Few Shot

Each point is one language direction, averaged equally across ten models. All three term types use the same directions.

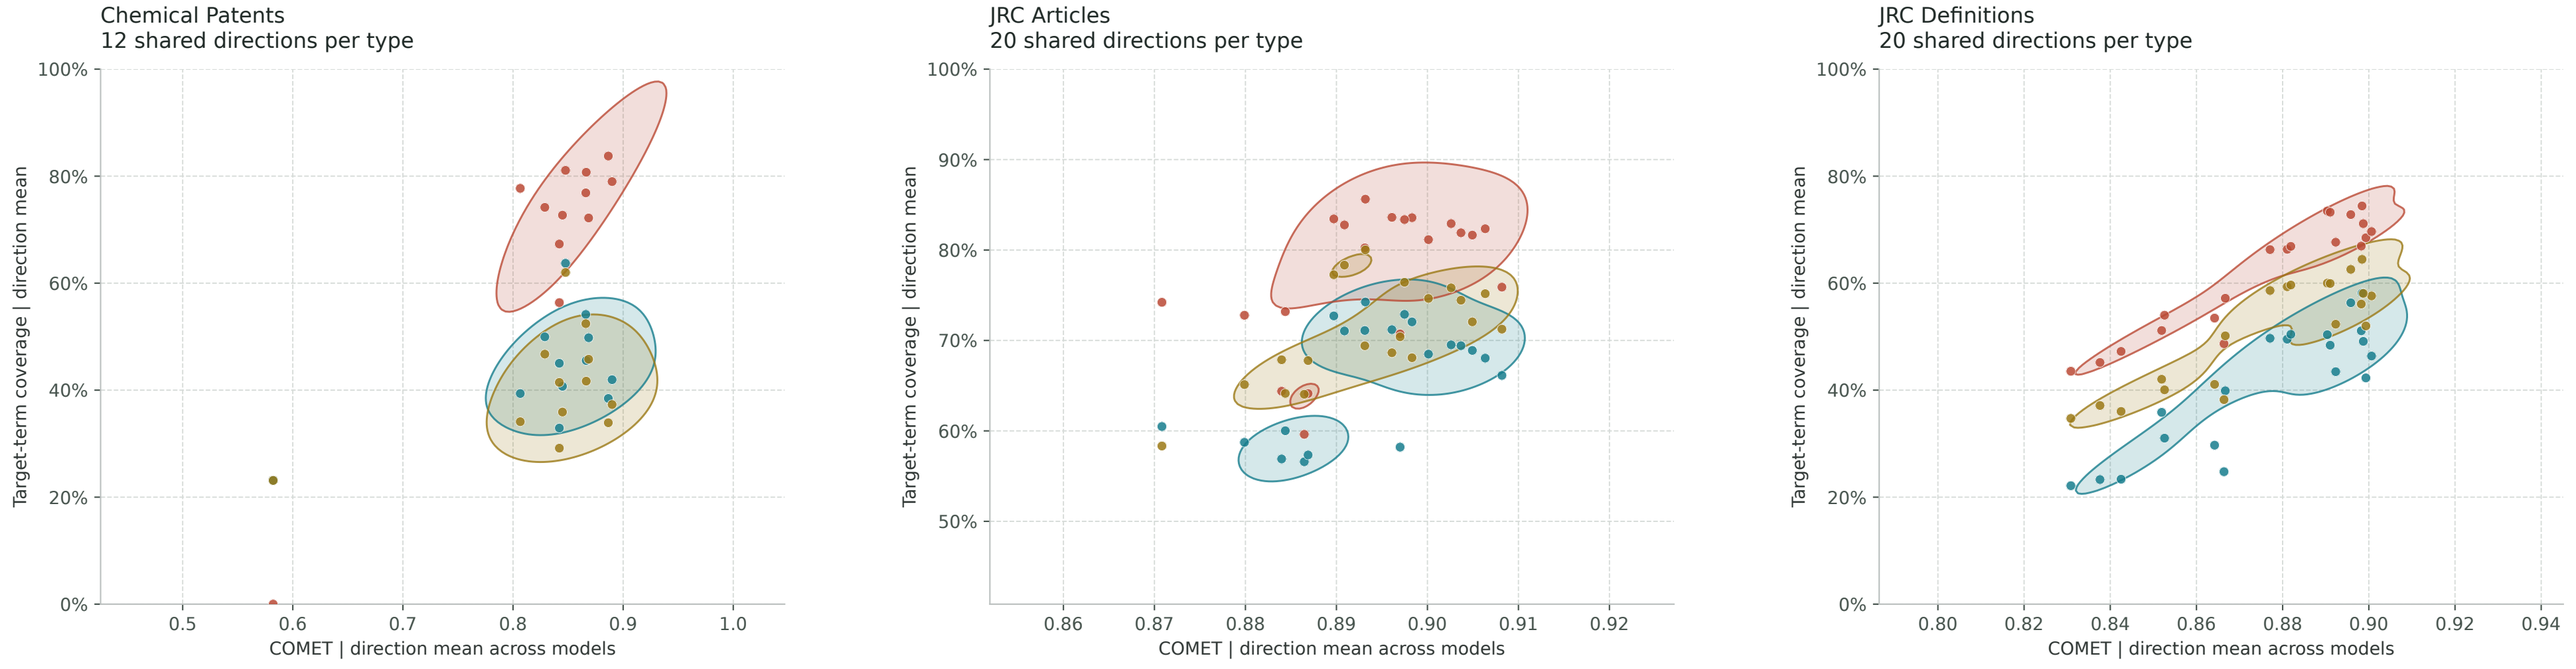

Shading: 50% fitted KDE mass | points-only fallback for sparse/degenerate data

All terms | Externally verified | LLM-enriched (refined)

Descriptive distributions of direction averages, not individual translations. KDE uses Scott bandwidth; regions are NOT confidence intervals.
Metric-specific scored rows may differ. Shared directions and models do not make observations independent; no significance inference.