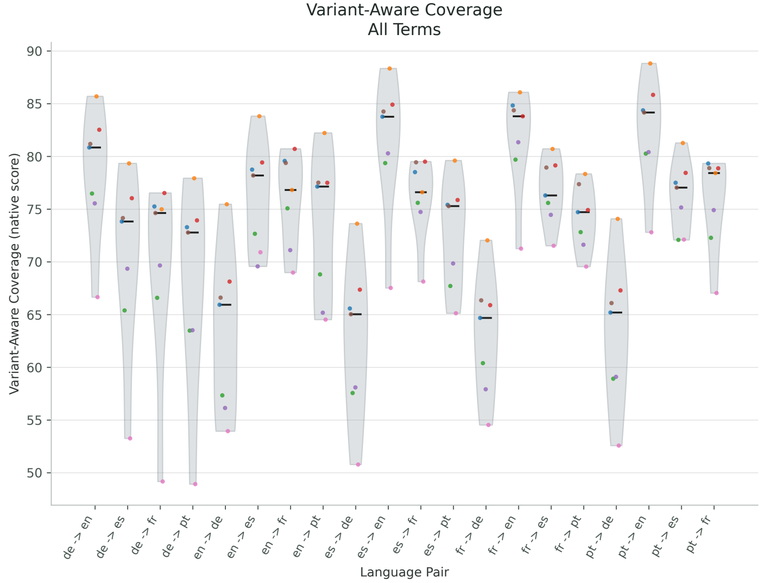
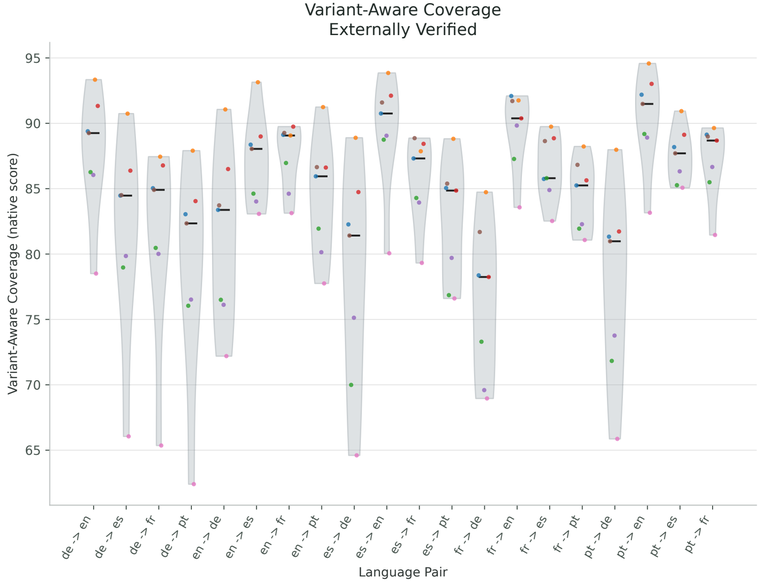
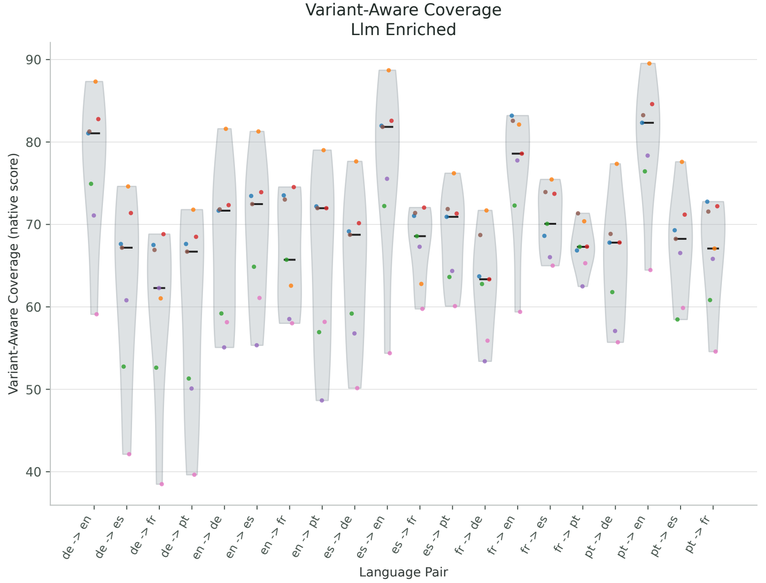
Performance Distributions by Language Pair | JRC Articles | luna | Few Shot







Gemma 4 31B · GPT-6 Luna · Granite 4.2 30B · Muse Glimmer 30B · Nemotron 3.5 Lightning · Qwen3.8 27B · Tiny Aya Global

One point = one model-direction score, NOT one translation. Equal model/direction contributions; native axes fit all observations.
Scott-bandwidth violin KDE; black line = median. Constant/sparse groups use points/median only. Partner sets and metric/term-group support may differ; no statistical inference.