Output Tokens vs Mixed Rank

JRC Articles | Zero Shot | luna | 5 default-thinking models

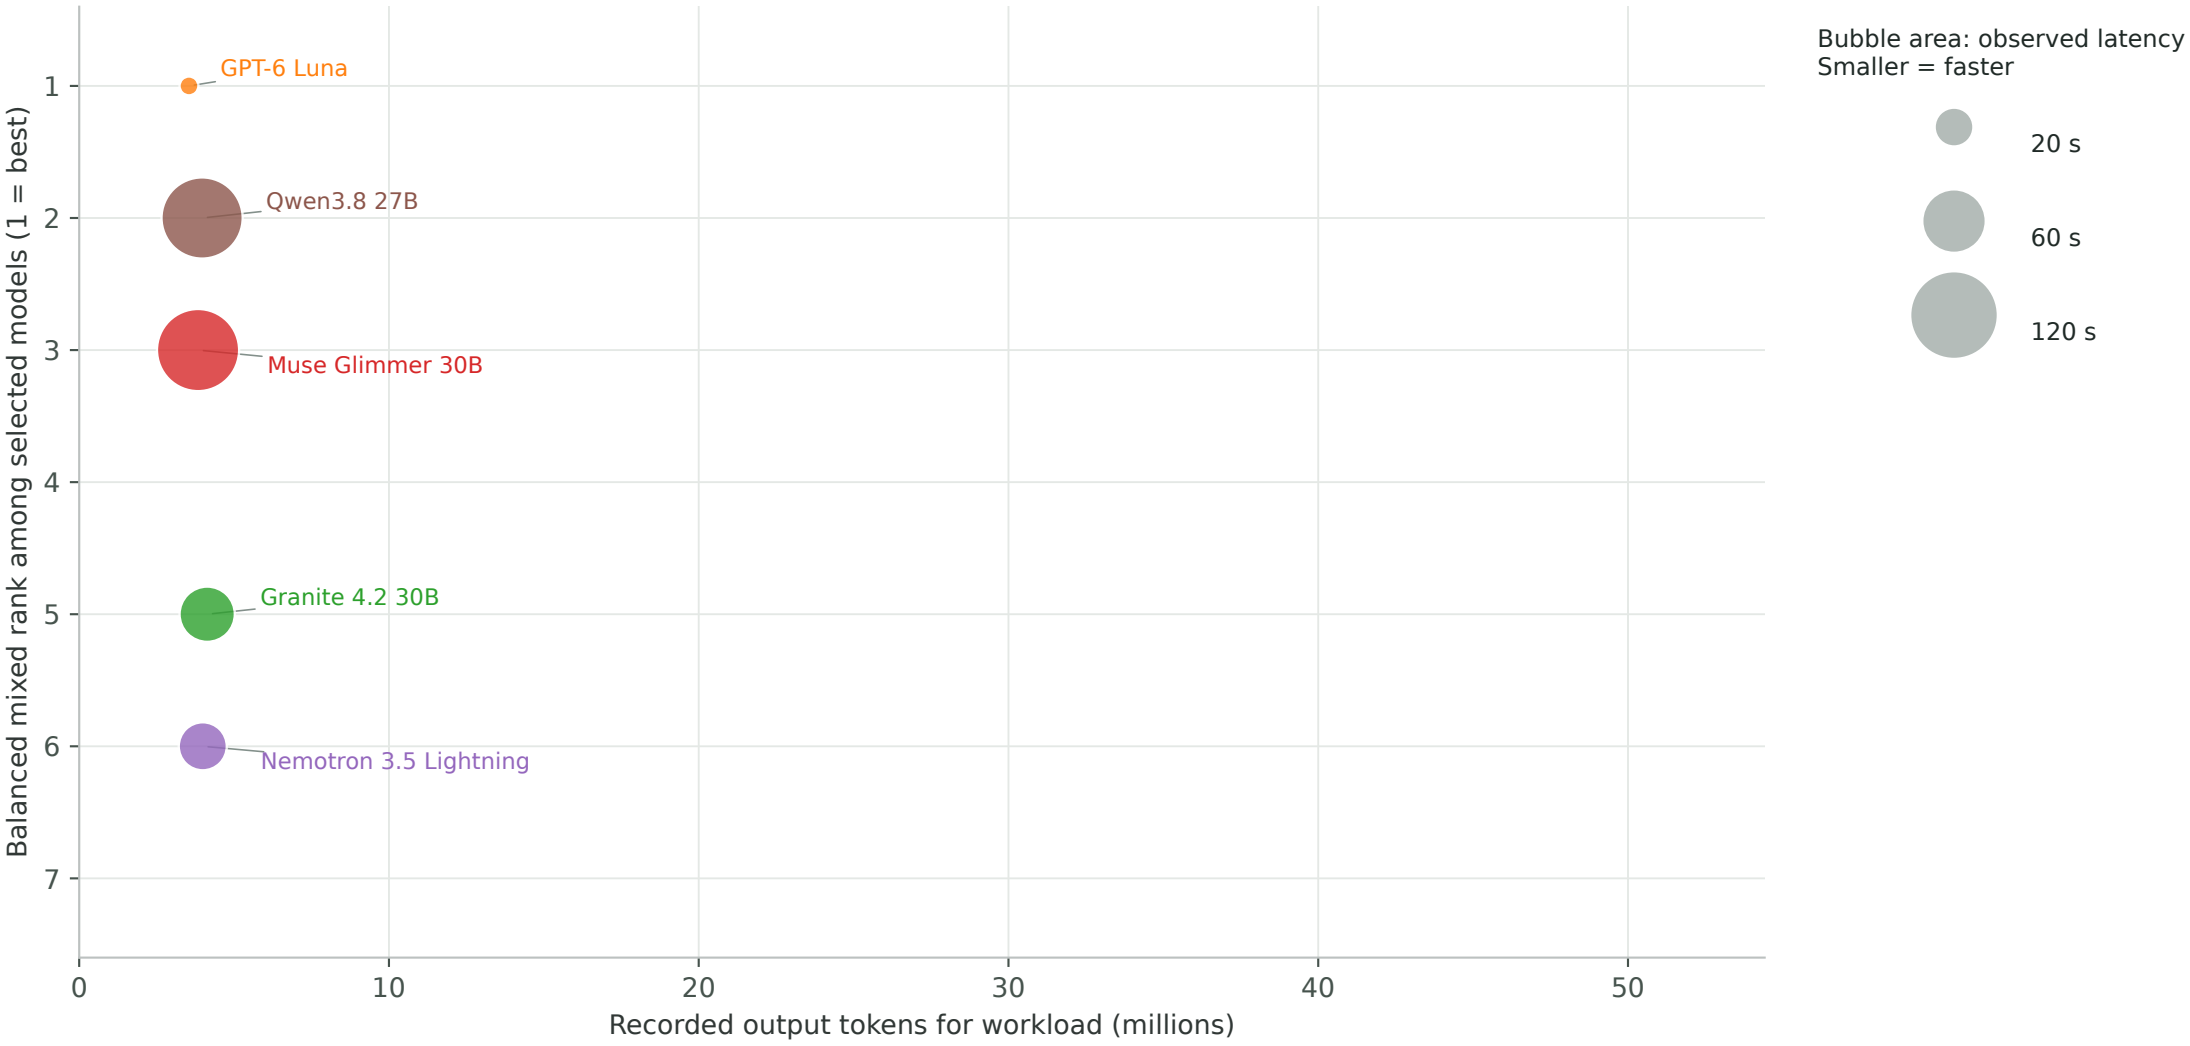

Original generation workload before analysis exclusions; not remeasured.

Area = 8 square points per latency second; identical token-axis and area scales across figures. Smaller bubbles mean lower observed latency.

Recorded output totals, NOT verified thinking-inclusive counts. Reported reasoning tokens are not added again.

Cohort: thinking enabled by default in official cards/templates; benchmark-run overrides are unverified.

Overall tokens sum six dataset/prompt strata; latency averages their means equally. Ranks remain among all selected models, not the subset.

Tokenizers, hardware, concurrency and workloads differ; not a row-matched or equal-budget comparison.

Excluded/missing: Gemma 4 31B, Tiny Aya Global.