Estimated API Cost vs Mixed Rank | Zoomed

JRC Definitions | Few Shot | luna | 5/7 models priced

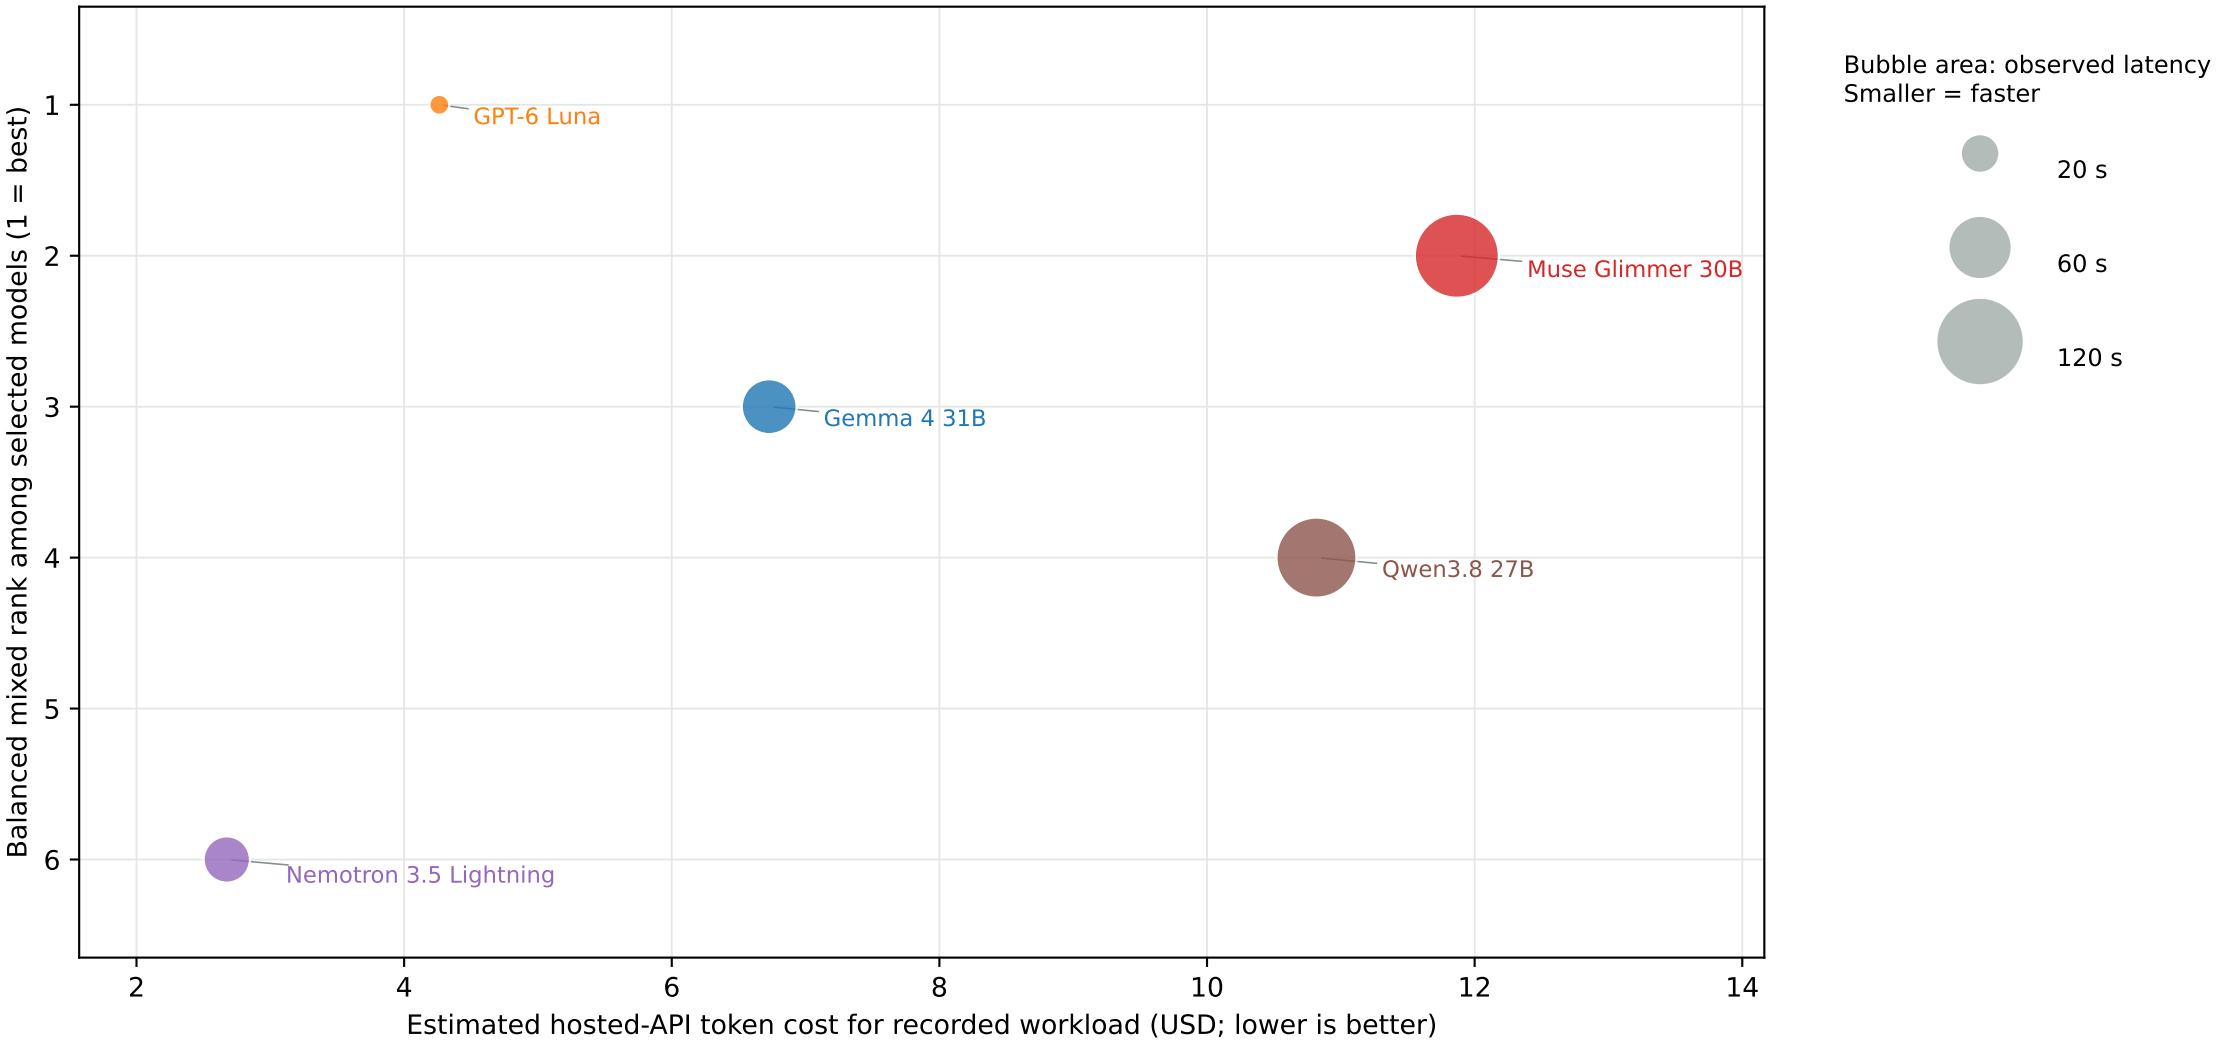

Original generation workload before analysis exclusions; not remeasured.

Upper-left with a smaller bubble is preferable. Area = 8 square points per latency second; identical area scale; axis ranges zoom to each panel without clipping.

Cost: published hosted-API reference repricing of recorded input + output tokens, no cache discount; not actual bills.

Latency: original local/API deployments, not the priced hosted offers. Hardware, concurrency and workloads differ; cost, timing and quality are not proven row-matched.

Overall cost sums six dataset/prompt strata; overall latency averages their means equally. Ranks are retained from all selected models, not recomputed for the priced subset.

Unpriced/missing (not zero): Granite 4.2 30B, Tiny Aya Global.