Do Metrics Agree About Prompting Changes?

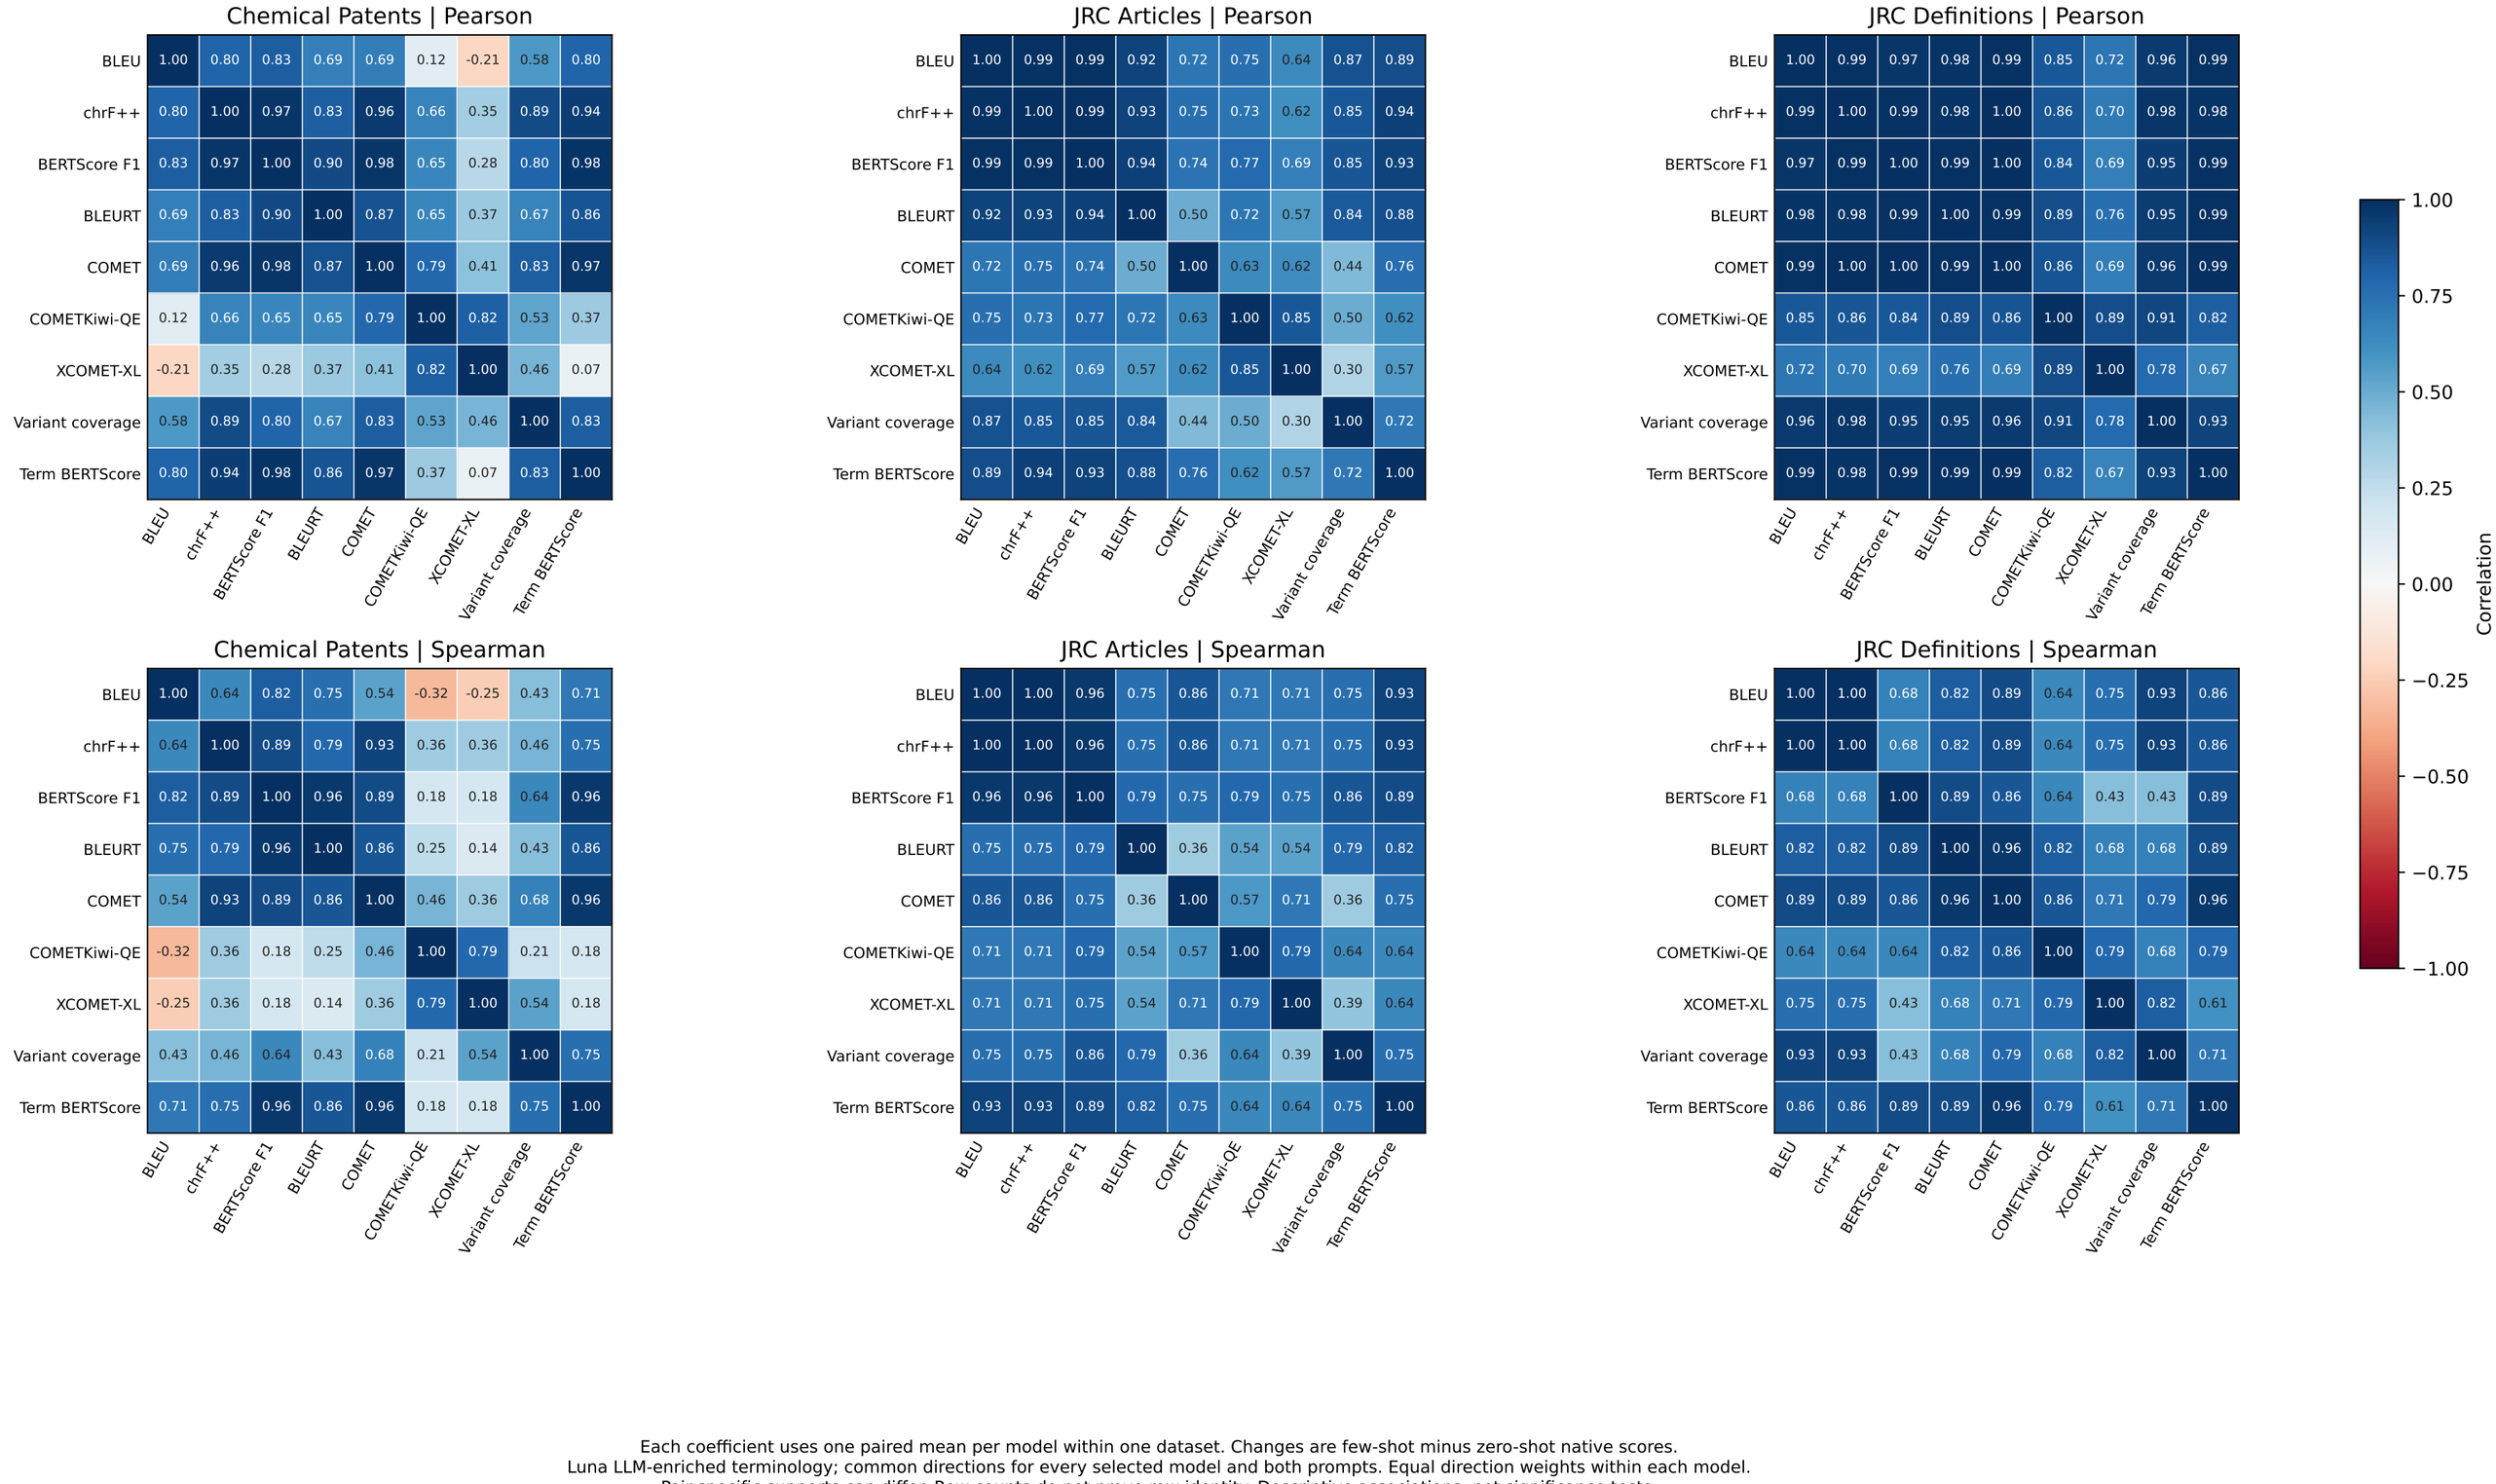

Each coefficient uses one paired mean per model within one dataset. Changes are few-shot minus zero-shot native scores.
Luna LLM-enriched terminology; common directions for every selected model and both prompts. Equal direction weights within each model.
Pair-specific supports can differ. Row counts do not prove row identity. Descriptive associations, not significance tests.